

```
In [ ]: import pandas as pd
```

```
In [ ]: df=pd.read_csv("Data/taxifare.csv",parse_dates=["pickup_datetime",
df.head()
```

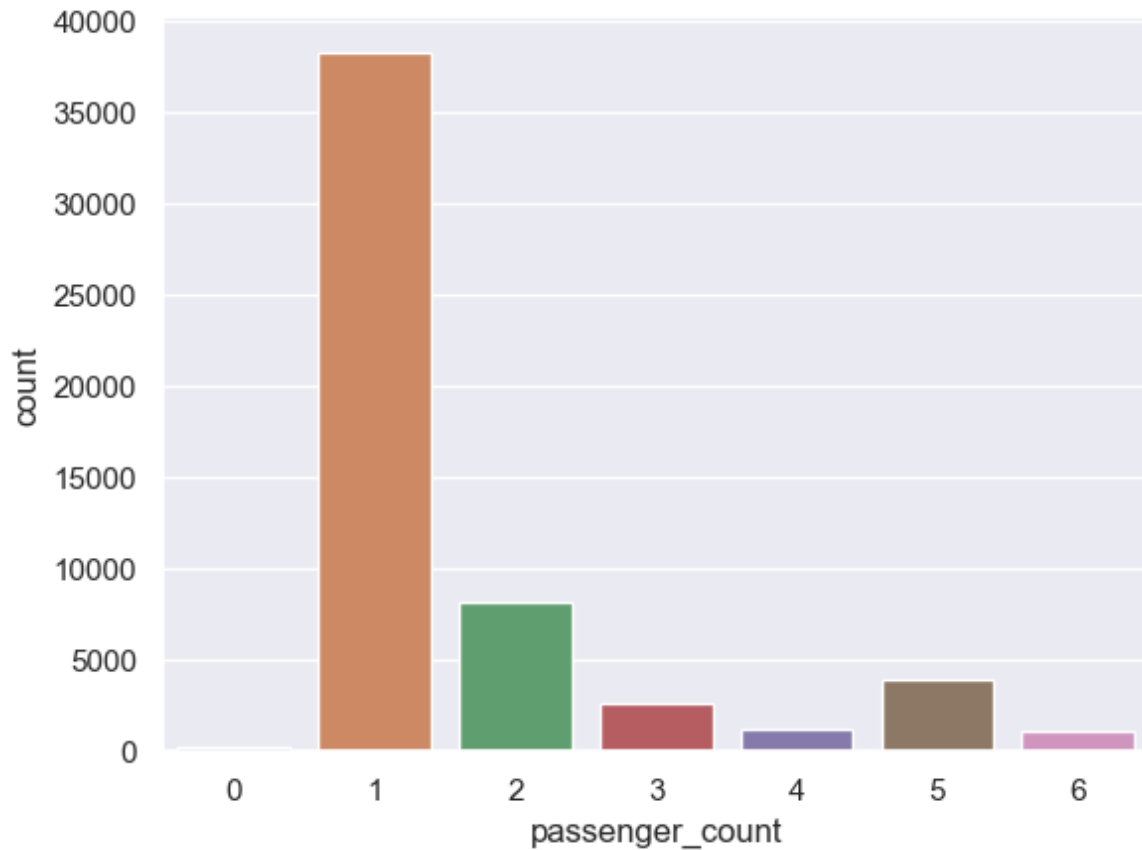
```
Out[ ]:
```

	key	fare_amount	pickup_datetime	pickup_longitude
<b>0</b>	2014-06-15 17:11:00.000000107	7.0	2014-06-15 17:11:00+00:00	-73.995420
<b>1</b>	2011-03-14 22:43:00.000000095	4.9	2011-03-14 22:43:00+00:00	-73.993552
<b>2</b>	2011-02-14 15:14:00.000000067	6.1	2011-02-14 15:14:00+00:00	-73.972380
<b>3</b>	2009-10-29 11:29:00.000000040	6.9	2009-10-29 11:29:00+00:00	-73.973703
<b>4</b>	2011-07-02 10:38:00.000000028	10.5	2011-07-02 10:38:00+00:00	-73.921262

The `parse_dates` argument tells Pandas to parse the `pickup_datetime` column as a date. This will allow you to use date and time operations on the data

```
In [ ]: %matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()
sns.countplot(x=df["passenger_count"])
```

```
Out[ ]: <Axes: xlabel='passenger_count', ylabel='count'>
```



Remove passenger with count  $\neq 1$

The `axis=1` parameter specifies that the columns should be dropped, as opposed to the rows. The `drop()` method takes a list of column names as its argument.

```
In [ ]: df = df[df["passenger_count"]==1]
df=df.drop(["key", "passenger_count"],axis=1)
df.head()
```

Out[ ]:

	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	d
0	7.0	2014-06-15 17:11:00+00:00	-73.995420	40.759662	
2	6.1	2011-02-14 15:14:00+00:00	-73.972380	40.749527	
4	10.5	2011-07-02 10:38:00+00:00	-73.921262	40.743615	
5	15.3	2011-12-09 20:03:00+00:00	-73.973500	40.792610	
8	7.7	2011-04-02 01:05:15+00:00	-73.979564	40.735405	

In [ ]: `df.shape`

Out[ ]: (38233, 6)

In [ ]: `corr_matrix=df.corr()  
corr_matrix["fare_amount"].sort_values(ascending=False)`

Out[ ]: fare\_amount 1.000000  
pickup\_datetime 0.115992  
dropoff\_longitude 0.020438  
pickup\_longitude 0.015742  
pickup\_latitude -0.015915  
dropoff\_latitude -0.021711  
Name: fare\_amount, dtype: float64

## Below

The code  $x=(\text{row}["\text{dropoff\_longitude}"]-\text{row}["\text{pickup\_longitude}"])/54.6$  and  $y=(\text{row}["\text{dropoff\_latitude}"]-\text{row}["\text{pickup\_latitude}"])/69.0$  will calculate the distance between the pickup and dropoff locations in miles and kilometers, respectively. The dropoff\_longitude and pickup\_longitude columns of the DataFrame df contain the longitudes of the pickup and dropoff locations, respectively. The row variable is a row from the DataFrame. The \* operator is used to multiply the two values together. The 54.6 and 69.0 values are the conversion factors from degrees to miles and kilometers, respectively.

In [ ]: `from math import sqrt  
for i,row in df.iterrows():  
 dt = row["pickup_datetime"]`

```

df.at[i,"day_of_week"]=dt.weekday()
df.at[i,"pickup_time"]=dt.hour
x=(row["dropoff_longitude"]-row["pickup_longitude"])*54.6
y=(row["dropoff_longitude"]-row["pickup_longitude"])*69.0
distance = sqrt(x**2+y**2)
df.at[i,"distance"]=distance
df.head()

```

Out[ ]:

	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	d
0	7.0	2014-06-15 17:11:00+00:00	-73.995420	40.759662	
2	6.1	2011-02-14 15:14:00+00:00	-73.972380	40.749527	
4	10.5	2011-07-02 10:38:00+00:00	-73.921262	40.743615	
5	15.3	2011-12-09 20:03:00+00:00	-73.973500	40.792610	
8	7.7	2011-04-02 01:05:15+00:00	-73.979564	40.735405	

## Drop unwanted columns

In [ ]: `df.drop(columns=["pickup_datetime","pickup_longitude","pickup_latitude"],inplace=True)`

In [ ]: `df.head()`

Out[ ]:

	fare_amount	day_of_week	pickup_time	distance
0	7.0	6.0	17.0	0.687462
2	6.1	0.0	15.0	1.606513
4	10.5	5.0	10.0	4.058166
5	15.3	4.0	20.0	3.296528
8	7.7	5.0	1.0	2.101014

In [ ]: `corr_matrix=df.corr()  
corr_matrix["fare_amount"].sort_values(ascending=False)`

```
Out[ ]: fare_amount    1.000000
        distance      0.045194
        day_of_week   0.009196
        pickup_time   -0.019722
        Name: fare_amount, dtype: float64
```

```
In [ ]: df.describe()
```

```
Out[ ]:
```

	fare_amount	day_of_week	pickup_time	distance
<b>count</b>	38233.000000	38233.000000	38233.000000	38233.000000
<b>mean</b>	11.214115	2.951534	13.387989	15.054776
<b>std</b>	9.703149	1.932809	6.446519	291.305494
<b>min</b>	-22.100000	0.000000	0.000000	0.000000
<b>25%</b>	6.000000	1.000000	9.000000	0.508316
<b>50%</b>	8.500000	3.000000	14.000000	1.087727
<b>75%</b>	12.500000	5.000000	19.000000	2.072325
<b>max</b>	256.000000	6.000000	23.000000	6514.524166

```
In [ ]: df = df[(df["distance"] > 1.0) & (df["distance"] < 10.0)]
df=df[(df["fare_amount"] > 0.0)&(df["fare_amount"] < 50.0)]
```

```
In [ ]: corr_matrix=df.corr()
corr_matrix["fare_amount"].sort_values(ascending=False)
```

```
Out[ ]: fare_amount    1.000000
        distance      0.728607
        day_of_week   0.003258
        pickup_time   -0.023683
        Name: fare_amount, dtype: float64
```

```
In [ ]: from sklearn.linear_model import LinearRegression
from sklearn.model_selection import cross_val_score

x= df.drop(["fare_amount"],axis=1)
y= df["fare_amount"]
```

## Model LinearRegressor

```
In [ ]: model= LinearRegression()
cross_val_score(model,x,y,cv=5).mean()
```

```
Out[ ]: 0.5295941284296068
```

## Model RandomForestRegressor

```
In [ ]: from sklearn.ensemble import RandomForestRegressor
model = RandomForestRegressor()
cross_val_score(model, x, y, cv=5).mean()
```

```
Out[ ]: 0.4611181767382222
```

## Model GradientBoostingRegressor

```
In [ ]: from sklearn.ensemble import GradientBoostingRegressor
model = GradientBoostingRegressor()
cross_val_score(model, x, y, cv=5).mean()
```

```
Out[ ]: 0.5402001030443373
```

## What is cv parameter?

cv is the number of folds to use for cross-validation. Cross-validation is a technique for evaluating the performance of a machine learning model on unseen data. It works by splitting the data into a training set and a test set. The model is trained on the training set and then evaluated on the test set. This process is repeated multiple times, with different splits of the data. The average of the scores from the multiple iterations is used as the final score for the model.

## Model Fitting

```
In [ ]: model.fit(x, y)
```

```
Out[ ]: ▾ GradientBoostingRegressor
GradientBoostingRegressor()
```

```
In [ ]: model.predict(pd.DataFrame({'day_of_week': [4], 'pickup_time':
```

```
Out[ ]: array([11.61950694])
```

```
In [ ]: model.predict(pd.DataFrame({'day_of_week': [5], 'pickup_time':
```

```
Out[ ]: array([11.27831854])
```

Regression models are supervised learning models that predict numeric outcomes such as the cost of a taxi ride. Prominent learning algorithms used for regression include the following:

Linear regression Models training data by fitting it to the equation of a line

Decision trees Use binary trees to predict an outcome by answering a series of yes-and-no questions

Random forests Use multiple independent decision trees to model the data and are resistant to overfitting

Gradient-boosting machines Use multiple dependent decision trees, each modeling the error in the output from the last

Support vector machines Take an entirely different approach to modeling data by adding dimensionality under the supposition that data that isn't linearly separable in the original problem space might be linearly separable in higher-dimensional space

Scikit provides convenient implementations of these and other learning algorithms in classes such as `LinearRegression`, `RandomForestRegressor`, and `GradientBoostingRegressor`

\$ k-fold cross-validation gives you more confidence in the  $R^2$  score than simply splitting the data once for training and testing. k-fold trains the model k times, each time with the dataset split differently.